

PSTAT 131 Homework Assignment 2

Marissa Santiago and Leticia Cruz

May 02, 2021

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.3    v purrr  0.3.4
## v tibble  3.1.0    v dplyr  1.0.5
## v tidyr   1.1.3    v stringr 1.4.0
## v readr   1.4.0    v forcats 0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(tree)
```

```
## Registered S3 method overwritten by 'tree':
##   method      from
##   print.tree cli
```

```
library(plyr)
```

```
## -----
## You have loaded plyr after dplyr - this is likely to cause problems.
## If you need functions from both plyr and dplyr, please load plyr first, then dplyr:
## library(plyr); library(dplyr)
```

```
## -----
```

```
##
## Attaching package: 'plyr'
```

```
## The following objects are masked from 'package:dplyr':
##
##   arrange, count, desc, failwith, id, mutate, rename, summarise,
##   summarize
```

```
## The following object is masked from 'package:purrr':
##
##   compact
```

```
library(dplyr)
library(class)
library(rpart)
library(maptree)
```

```
## Loading required package: cluster
```

```
library(ROCR)
library(reshape2)
```

```
##
## Attaching package: 'reshape2'

## The following object is masked from 'package:tidyr':
##
## smiths
```

```
spam <- read_table2("spambase.tab", guess_max=2000)
```

```
##
## -- Column specification -----
## cols(
##   .default = col_double()
## )
## i Use `spec()` for the full column specifications.
```

```
spam <- spam %>%
  mutate(y = factor(y, levels=c(0,1), labels=c("good", "spam"))) %>%
  mutate_at(.vars=vars(-y), .funs=scale)
```

```
calc_error_rate <- function(predicted.value, true.value){
  return(mean(true.value!=predicted.value))
}
```

```
records = matrix(NA, nrow=3, ncol=2)
colnames(records) <- c("train.error", "test.error")
rownames(records) <- c("knn", "tree", "logistic")
```

Training/test sets: Split randomly the data set in a train and a test set:

```
set.seed(1)
test.indices = sample(1:nrow(spam), 1000)
spam.train=spam[-test.indices,]
spam.test=spam[test.indices,]
```

Folds for cv

```

nfold = 10
set.seed(1)
folds = seq.int(nrow(spam.train)) %>%
  cut(breaks = nfold, labels=FALSE) %>%
  sample

```

Problem 1

```

set.seed(1)
do.chunk <- function(chunkid, folddef, Xdat, Ydat, k){
  train = (folddef!=chunkid)
  Xtr = Xdat[train,]
  Ytr = Ydat[train]
  Xvl = Xdat[!train,]
  Yvl = Ydat[!train]
  ## get classifications for current training chunks
  predYtr = knn(train = Xtr, test = Xtr, cl = Ytr, k = k)
  ## get classifications for current test chunk
  predYvl = knn(train = Xtr, test = Xvl, cl = Ytr, k = k)

  data.frame(fold = chunkid,
             train.error = calc_error_rate(predYtr, Ytr),
             val.error = calc_error_rate(predYvl, Yvl))
}

```

```

#check for missing values
sum(is.na(spam.train))

```

```
## [1] 0
```

```
sum(is.na(spam.test))
```

```
## [1] 0
```

```

YTrain = spam.train$y %>% na.omit()
XTrain = spam.train %>% select(-y) %>% na.omit()
YTest = spam.test$y %>% na.omit()
XTest = spam.test %>% select(-y) %>% na.omit()

error.folds <- NULL
set.seed(1)

kvec = c(1, seq(10, 50, length.out=5))
for (i in kvec){
  tmp <- ldply(1:nfold, do.chunk,
             folddef = folds, Xdat = XTrain,
             Ydat = YTrain, k = i)
  tmp$neighbors <- i
  error.folds <- rbind(error.folds,tmp)
}
error.folds

```

##	fold	train.error	val.error	neighbors
## 1	1	0.0006173	0.11080	1
## 2	2	0.0000000	0.11944	1
## 3	3	0.0006171	0.08056	1
## 4	4	0.0000000	0.08056	1
## 5	5	0.0006171	0.10833	1
## 6	6	0.0006171	0.11111	1
## 7	7	0.0003085	0.07778	1
## 8	8	0.0000000	0.11667	1
## 9	9	0.0003085	0.10000	1
## 10	10	0.0003085	0.13056	1
## 11	1	0.0824074	0.08864	10
## 12	2	0.0823820	0.11111	10
## 13	3	0.0805307	0.08889	10
## 14	4	0.0774452	0.10000	10
## 15	5	0.0755940	0.09722	10
## 16	6	0.0762110	0.10278	10
## 17	7	0.0805307	0.05833	10
## 18	8	0.0789880	0.09444	10
## 19	9	0.0759025	0.11111	10
## 20	10	0.0786794	0.11389	10
## 21	1	0.0919753	0.09418	20
## 22	2	0.0944153	0.11944	20
## 23	3	0.0956495	0.08056	20
## 24	4	0.0934897	0.08889	20
## 25	5	0.0888615	0.12500	20
## 26	6	0.0882444	0.11111	20
## 27	7	0.0965751	0.06944	20
## 28	8	0.0907127	0.10556	20
## 29	9	0.0931811	0.12778	20
## 30	10	0.0910213	0.10000	20
## 31	1	0.0993827	0.10249	30
## 32	2	0.1024375	0.12500	30
## 33	3	0.1052144	0.10000	30
## 34	4	0.1030546	0.10556	30
## 35	5	0.0993521	0.11667	30
## 36	6	0.0984264	0.10833	30
## 37	7	0.1033632	0.07778	30
## 38	8	0.0971922	0.12778	30
## 39	9	0.1012033	0.11944	30
## 40	10	0.0990435	0.10833	30
## 41	1	0.1055556	0.11357	40
## 42	2	0.1058315	0.11667	40
## 43	3	0.1104597	0.11111	40
## 44	4	0.1052144	0.10556	40
## 45	5	0.1101512	0.12222	40
## 46	6	0.1073743	0.12222	40
## 47	7	0.1098426	0.08056	40
## 48	8	0.1021290	0.13056	40
## 49	9	0.1098426	0.12500	40
## 50	10	0.1033632	0.09722	40
## 51	1	0.1111111	0.11080	50
## 52	2	0.1129281	0.12222	50
## 53	3	0.1135452	0.11389	50

```
## 54    4  0.1110768  0.11111    50
## 55    5  0.1104597  0.12222    50
## 56    6  0.1116939  0.11944    50
## 57    7  0.1141623  0.08056    50
## 58    8  0.1110768  0.14722    50
## 59    9  0.1082999  0.12222    50
## 60   10  0.1089170  0.10556    50
```

```
#Transform the format of error.folds for further convenience
errors = melt(error.folds, id.vars=c('fold', 'neighbors'), value.name='error')

val.error.means = errors %>%
  # Select all rows of validation errors
  filter(variable=='val.error') %>%
  # Group the selected data frame by neighbors
  group_by(neighbors, variable) %>%
  # Calculate CV error rate for each k
  summarise_each(funs(mean), error) %>%
  # Remove existing group
  ungroup() %>%
  filter(error==min(error))
```

```
## Warning: `summarise_each()` was deprecated in dplyr 0.7.0.
## Please use `across()` instead.
```

```
## Warning: `funs()` was deprecated in dplyr 0.8.0.
## Please use a list of either functions or lambdas:
```

```
##
## # Simple named list:
## list(mean = mean, median = median)
##
## # Auto named with `tibble::lst()`:
## tibble::lst(mean, median)
##
## # Using lambdas
## list(~ mean(., trim = .2), ~ median(., na.rm = TRUE))
```

```
val.error.means
```

```
## # A tibble: 1 x 3
##   neighbors variable    error
##       <dbl> <fct>      <dbl>
## 1         10 val.error 0.0966
```

```
best.kfold = max(val.error.means$neighbors)
best.kfold
```

```
## [1] 10
```

Problem 2

```

set.seed(1)
#training error rate
pred.YTrain = knn(train = XTrain, test = XTrain, cl = YTrain, k = best.kfold)
train_error = calc_error_rate(pred.YTrain,YTrain)

#test error rate
pred.YTest = knn( train = XTrain, test = XTest, cl = YTrain, k = best.kfold)
test_error = calc_error_rate(pred.YTest, YTest)

records[1,1]=train_error
records[1,2]=test_error
records

```

```

##          train.error test.error
## knn          0.07803      0.102
## tree              NA         NA
## logistic        NA         NA

```

Problem 3

```

spamtrees = tree(y ~ ., data = spam.train,
                 control = tree.control(nrow(spam.train), minsize = 5, mindev = 1e-5))
summary(spamtrees)

```

```

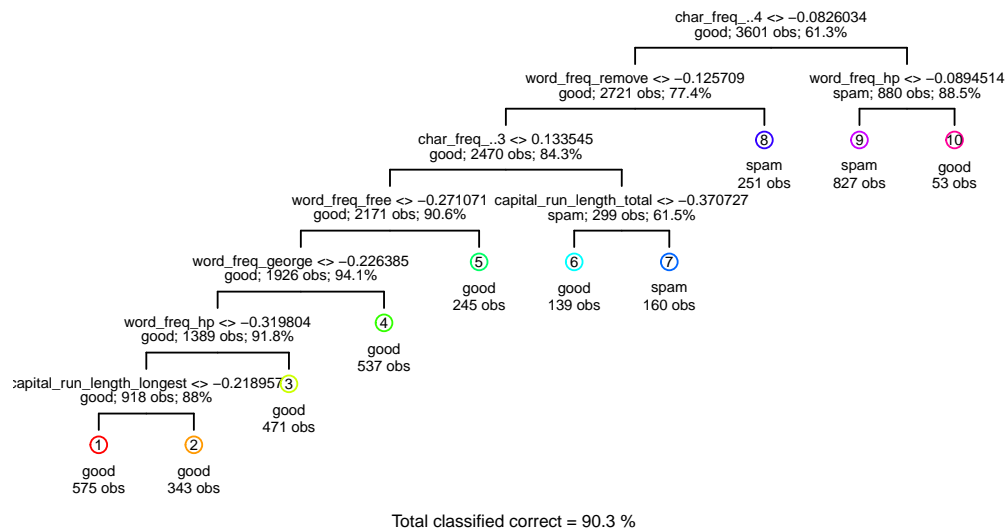
##
## Classification tree:
## tree(formula = y ~ ., data = spam.train, control = tree.control(nrow(spam.train),
##   minsize = 5, mindev = 1e-05))
## Variables actually used in tree construction:
## [1] "char_freq..4"          "word_freq_remove"
## [3] "char_freq..3"          "word_freq_free"
## [5] "word_freq_george"      "word_freq_hp"
## [7] "capital_run_length_longest" "word_freq_receive"
## [9] "word_freq_credit"      "capital_run_length_average"
## [11] "word_freq_your"        "word_freq_mail"
## [13] "word_freq_re"          "word_freq_our"
## [15] "word_freq_you"         "capital_run_length_total"
## [17] "word_freq_make"        "word_freq_all"
## [19] "word_freq_internet"    "word_freq_email"
## [21] "word_freq_project"     "word_freq_money"
## [23] "word_freq_1999"        "word_freq_will"
## [25] "char_freq..1"          "word_freq_order"
## [27] "char_freq.."           "word_freq_data"
## [29] "word_freq_over"        "word_freq_meeting"
## [31] "word_freq_650"         "word_freq_edu"
## [33] "word_freq_address"     "word_freq_business"
## Number of terminal nodes: 149
## Residual mean deviance: 0.0457 = 158 / 3450
## Misclassification error rate: 0.0136 = 49 / 3601

```

There is a total of 149 leaf nodes in this tree and there are 49 training observations that are misclassified.

Problem 4

```
prune <- prune.tree(spamtree,best = 10)
draw.tree(prune, nodeinfo=TRUE, cex = 0.5)
```



Problem 5

```
set.seed(1)
cv = cv.tree(spamtree, rand = folds, FUN = prune.misclass, K = 10)
cv
```

```
## $size
## [1] 149 106 102 99 76 73 63 59 52 46 41 38 35 24 22 16 15 14 13
## [20] 9 8 7 6 5 4 3 2 1
##
## $dev
## [1] 353 353 346 346 346 346 346 346 346 346 346 346 342 342 342 342
## [16] 344 344 343 343 347 354 352 352 410 463 516 715 1393
##
## $k
## [1] -Inf 0.0000 0.5000 0.6667 1.0000 1.3333 1.5000 1.7500
```

```
## [9] 2.0000 2.5000 2.8000 3.0000 3.6667 4.0000 4.5000 5.1667
## [17] 6.0000 7.0000 8.0000 9.7500 11.0000 12.0000 17.0000 45.0000
## [25] 53.0000 69.0000 199.0000 678.0000
##
## $method
## [1] "misclass"
##
## attr("class")
## [1] "prune"          "tree.sequence"
```

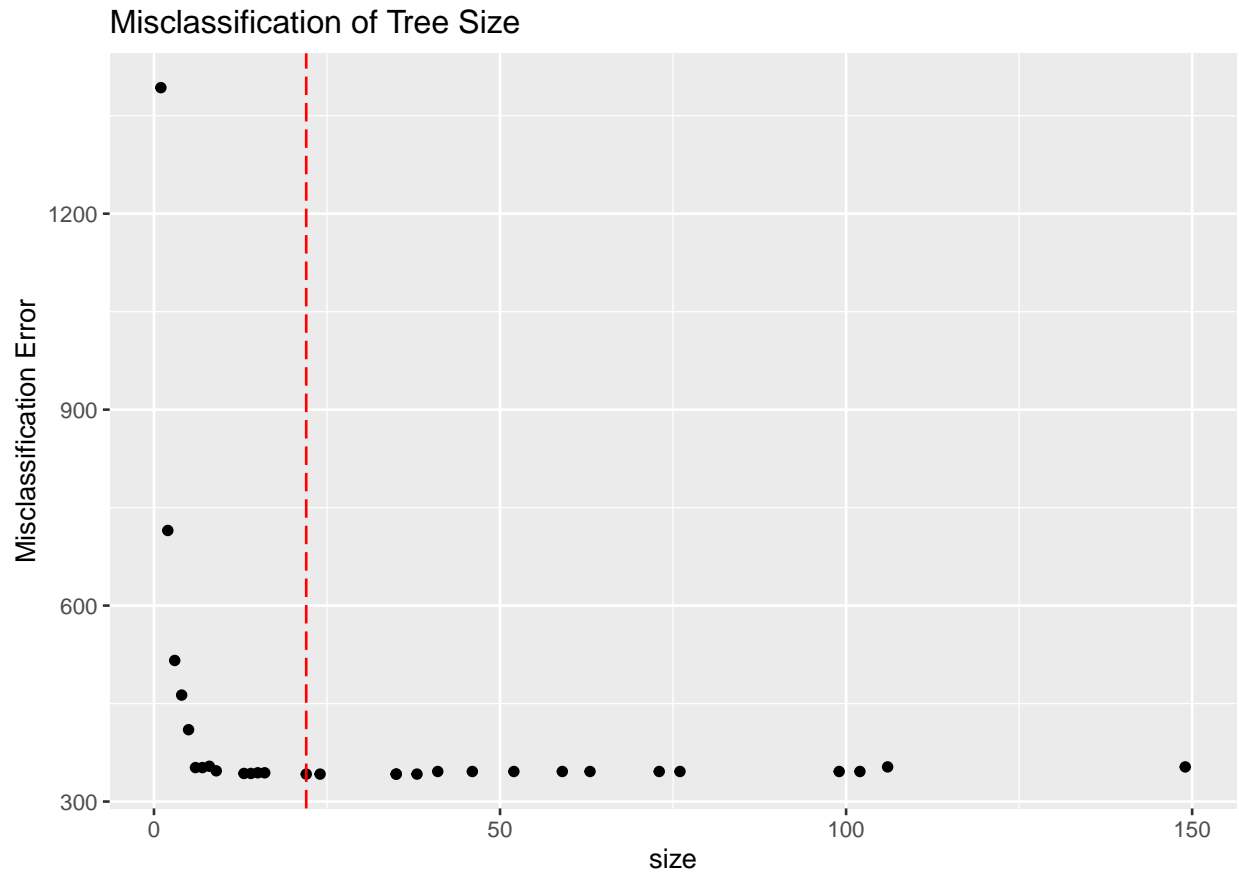
```
tree_data <- data.frame("x"=cv$size, "y"=cv$dev)
tree_data
```

```
##      x      y
## 1 149 353
## 2 106 353
## 3 102 346
## 4  99 346
## 5  76 346
## 6  73 346
## 7  63 346
## 8  59 346
## 9  52 346
## 10 46 346
## 11 41 346
## 12 38 342
## 13 35 342
## 14 24 342
## 15 22 342
## 16 16 344
## 17 15 344
## 18 14 343
## 19 13 343
## 20  9 347
## 21  8 354
## 22  7 352
## 23  6 352
## 24  5 410
## 25  4 463
## 26  3 516
## 27  2 715
## 28  1 1393
```

```
#best size
best.size.cv = min(cv$size[cv$dev==min(cv$dev)])
best.size.cv
```

```
## [1] 22
```

```
ggplot(tree_data,aes(x,y)) +geom_point() + geom_point(data=tree_data[13,],aes(x,y)) + geom_vline(xinter
ggtitle("Misclassification of Tree Size")+xlab("size")+ylab("Misclassification Error")
```

The optimal tree size is 22.

Problem 6

```
set.seed(1)
#training error
spamtree.pruned = prune.misclass(spamtree, best = best.size.cv)

pred.train = predict(spamtree.pruned, spam.train, type = "class")
#testing error
pred.test = predict(spamtree.pruned, spam.test, type = "class")

train.error = calc_error_rate(pred.train, YTrain)
test.error = calc_error_rate(pred.test, YTest)

records[2,2] = test.error
records[2,1] = train.error
records
```

```
##          train.error test.error
## knn          0.07803      0.102
## tree          0.06054      0.091
## logistic          NA          NA
```

Logistics Regression

Problem 7

7a

Given,

$$p(z) = \frac{e^z}{1 + e^z}$$

$$p(1 + e^z) = e^z$$

$$p + pe^z = e^z$$

$$p = e^z - pe^z$$

$$p = e^z(1 - p)$$

$$\frac{p}{(1 - p)} = e^z$$

$$e^z = \frac{p}{(1 - p)}$$

$$z = \ln\left(\frac{p}{1 - p}\right)$$

7b

#

$$p = \frac{e^{\beta_0 + \beta_1 x_1}}{1 + e^{\beta_0 + \beta_1 x_1}}$$

When $x_1 = x_1 + 2$:

#

$$p = \frac{e^{\beta_0 + \beta_1(x_1 + 2)}}{1 + e^{\beta_0 + \beta_1(x_1 + 2)}} = \frac{e^{\beta_0 + \beta_1 x_1 + 2\beta_1}}{1 + e^{\beta_0 + \beta_1 x_1 + 2\beta_1}} = \frac{e^{\beta_0 + \beta_1 x_1} e^{2\beta_1}}{1 + e^{\beta_0 + \beta_1 x_1} e^{2\beta_1}}$$

As x increases by 2, the odds is multiplied by $e^{2\beta_1}$.

#

$$\lim_{x \rightarrow \infty} p = 0$$

Also, as x goes to infinity, the numerator becomes smaller and the denominator becomes bigger. Therefore, the probability gets closer to 0.

#

$$\lim_{x \rightarrow -\infty} p = 1$$

As x goes to negative infinity, the probability goes to 1.

Problem 8

```
set.seed(1)
#fit logistic regression
glm.fit = glm(y~.,data=spam.train, family=binomial)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(glm.fit)
```

```
##
## Call:
## glm(formula = y ~ ., family = binomial, data = spam.train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.812  -0.198   0.000   0.119   5.551
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -1.64e+01   1.98e+02  -0.08  0.93399
## word_freq_make    -1.14e-01   7.73e-02  -1.48  0.13990
## word_freq_address -1.58e-01   9.42e-02  -1.68  0.09270 .
## word_freq_all      8.79e-02   6.13e-02   1.43  0.15175
## word_freq_3d       3.66e+00   2.41e+00   1.52  0.12789
## word_freq_our       4.83e-01   8.03e-02   6.02  1.7e-09 ***
## word_freq_over      2.91e-01   8.18e-02   3.56  0.00037 ***
## word_freq_remove    7.89e-01   1.30e-01   6.08  1.2e-09 ***
## word_freq_internet  1.92e-01   6.55e-02   2.94  0.00331 **
## word_freq_order     1.57e-01   8.67e-02   1.81  0.07053 .
## word_freq_mail      5.91e-02   4.78e-02   1.24  0.21638
## word_freq_receive  -4.91e-02   6.57e-02  -0.75  0.45412
## word_freq_will     -1.25e-01   7.34e-02  -1.70  0.08919 .
## word_freq_people   -2.35e-03   8.05e-02  -0.03  0.97673
## word_freq_report    1.46e-02   5.20e-02   0.28  0.77940
## word_freq_addresses 3.00e-01   1.83e-01   1.64  0.10177
## word_freq_free      8.92e-01   1.33e-01   6.70  2.1e-11 ***
## word_freq_business 3.53e-01   1.03e-01   3.42  0.00063 ***
## word_freq_email     9.84e-02   6.69e-02   1.47  0.14153
## word_freq_you       1.28e-01   6.95e-02   1.84  0.06519 .
## word_freq_credit    5.14e-01   3.12e-01   1.65  0.09904 .
## word_freq_your      2.61e-01   6.92e-02   3.77  0.00017 ***
## word_freq_font      3.17e-01   2.30e-01   1.38  0.16857
## word_freq_000       8.18e-01   1.85e-01   4.42  9.9e-06 ***
## word_freq_money     1.99e-01   7.42e-02   2.69  0.00721 **
## word_freq_hp       -3.36e+00   6.06e-01  -5.54  3.0e-08 ***
## word_freq_hpl      -7.02e-01   3.93e-01  -1.79  0.07410 .
## word_freq_george   -4.13e+01   8.45e+00  -4.88  1.0e-06 ***
## word_freq_650       2.67e-01   1.87e-01   1.43  0.15378
## word_freq_lab      -1.23e+00   8.37e-01  -1.46  0.14319
## word_freq_labs     -1.80e-01   1.73e-01  -1.04  0.29990
## word_freq_telnet   -4.72e-02   1.50e-01  -0.31  0.75357
```

```
## word_freq_857          -2.52e+01  1.38e+03  -0.02  0.98547
## word_freq_data        -5.96e-01  2.19e-01  -2.72  0.00648 **
## word_freq_415         3.96e-01  5.80e-01   0.68  0.49460
## word_freq_85          -1.08e+00  4.54e-01  -2.38  0.01710 *
## word_freq_technology   2.58e-01  1.43e-01   1.80  0.07114 .
## word_freq_1999        4.45e-02  8.14e-02   0.55  0.58481
## word_freq_parts       3.71e-01  2.13e-01   1.74  0.08205 .
## word_freq_pm          -2.81e-01  1.95e-01  -1.44  0.15065
## word_freq_direct      -1.12e-01  1.33e-01  -0.84  0.39891
## word_freq_cs          -1.68e+01  9.60e+00  -1.75  0.08067 .
## word_freq_meeting     -2.45e+00  8.55e-01  -2.87  0.00414 **
## word_freq_original    -1.57e-01  1.62e-01  -0.97  0.33165
## word_freq_project     -1.14e+00  3.94e-01  -2.88  0.00397 **
## word_freq_re          -7.10e-01  1.54e-01  -4.60  4.2e-06 ***
## word_freq_edu         -1.21e+00  2.59e-01  -4.68  2.9e-06 ***
## word_freq_table       -1.10e-01  1.42e-01  -0.78  0.43734
## word_freq_conference  -1.31e+00  5.56e-01  -2.35  0.01894 *
## char_freq_            -4.15e-01  1.55e-01  -2.68  0.00739 **
## char_freq_..1        -3.96e-02  8.33e-02  -0.48  0.63461
## char_freq_..2        -6.59e-02  1.16e-01  -0.57  0.56942
## char_freq_..3         1.97e-01  5.54e-02   3.56  0.00037 ***
## char_freq_..4         1.08e+00  1.83e-01   5.89  3.8e-09 ***
## char_freq_..5         1.22e+00  4.99e-01   2.45  0.01443 *
## capital_run_length_average 3.17e-01  6.62e-01   0.48  0.63224
## capital_run_length_longest 1.79e+00  5.60e-01   3.20  0.00139 **
## capital_run_length_total  7.14e-01  1.53e-01   4.68  2.9e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 4806.0  on 3600  degrees of freedom
## Residual deviance: 1413.2  on 3543  degrees of freedom
## AIC: 1529
##
## Number of Fisher Scoring iterations: 22
```

```
#test
prob.train=predict(glm.fit,type="response")
prob.test=predict(glm.fit,newdata=spam.test,type="response")

#save the predicted labels
spamtrain = spam.train%>%
  mutate(predspamtrain=as.factor(ifelse(prob.train<=0.5,"good","spam")))
spamttest = spam.test%>%
  mutate(predspamttest=as.factor(ifelse(prob.test<=0.5,"good","spam")))
d<-calc_error_rate(spamtrain$predspamtrain,YTrain)
e<-calc_error_rate(spamttest$predspamttest,YTest)
records[3,1]=d
records[3,2]=e
records
```

```
##          train.error test.error
## knn      0.07803      0.102
```

## tree	0.06054	0.091
## logistic	0.06804	0.086

The method with the lowest misclassification error is decision tree method.

Problem 9

We take “positive” here to mean “spam.” We would be more concerned with the false positive rate being large. In this case our model would not be correctly classifying the email as “spam” which means that all those spam emails will end up in our inbox; leading to a mixture of spam emails with legitimate emails. With a low true positive rate, the spam email would get filtered into where it belongs so we wouldn’t be too concerned with this. The user can always delete spam from their regular inbox , but cannot easily recover or notice good emails being placed in the spam folder.